

# Exploratory Data Analysis and Prediction of IBM Employees.

DataSet Source : <https://www.kaggle.com/datasets/rohitsahoo/employee>

This dataset have 1058 Rows and 36 Columns.

## About Dataset

### Education

- 'Below College'
- 'College'
- 'Bachelor'
- 'Master'
- 'Doctor'

### EnvironmentSatisfaction

- 'Low'
- 'Medium'
- 'High'
- 'Very High'

### JobInvolvement

- 'Low'
- 'Medium'
- 'High'
- 'Very High'

### JobSatisfaction

- 'Low'
- 'Medium'
- 'High'
- 'Very High'

### PerformanceRating

- 'Low'
- 'Good'
- 'Excellent'

- 'Outstanding'

## RelationshipSatisfaction

- 'Low'
- 'Medium'
- 'High'
- 'Very High'

## WorkLifeBalance

- 'Bad'
- 'Good'
- 'Better'
- 'Best'

## This notebook is structured as follows:

- Exploratory Data Analysis: In this section, we explore the dataset by taking a look at the feature distributions, how correlated one feature is to the other and create some Seaborn and Plotly visualisations
- Feature Engineering and Categorical Encoding: Conduct some feature engineering as well as encode all our categorical features into dummy variables
- Implementing Machine Learning models: We implement a Random Forest and a Gradient Boosted Model after which we look at feature importances from these respective models

Let's Go.

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
sns.set()
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

# EDA and Prediction of train.csv Dataset.

## 1. Exploratory Data Analysis.

```
In [3]: data=pd.read_csv('train.csv')
```

```
In [75]: #Top 5 Records.

data.head()
```

Out[75]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education
0	41	1	Travel_Rarely	1102	Sales	1	2	Life Sci
1	49	0	Travel_Frequently	279	Research & Development	8	1	Life Sci
2	37	1	Travel_Rarely	1373	Research & Development	2	2	
3	33	0	Travel_Frequently	1392	Research & Development	3	4	Life Sci
4	27	0	Travel_Rarely	591	Research & Development	2	1	Me

5 rows × 36 columns



## Data quality checks

To look for any null values, we can just invoke the `isnull` call as follows

In [5]: `data.isnull().any()`

```
Out[5]: Age False
Attrition False
BusinessTravel False
DailyRate False
Department False
DistanceFromHome False
Education False
EducationField False
EmployeeCount False
EmployeeNumber False
EnvironmentSatisfaction False
Gender False
HourlyRate False
JobInvolvement False
JobLevel False
JobRole False
JobSatisfaction False
MaritalStatus False
MonthlyIncome False
MonthlyRate False
NumCompaniesWorked False
Over18 False
OverTime False
PercentSalaryHike False
PerformanceRating False
RelationshipSatisfaction False
StandardHours False
StockOptionLevel False
TotalWorkingYears False
TrainingTimesLastYear False
WorkLifeBalance False
YearsAtCompany False
YearsInCurrentRole False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool
```

```
In [6]: data.dtypes
```

```
Out[6]: Age                int64
Attrition                int64
BusinessTravel           object
DailyRate                int64
Department              object
DistanceFromHome         int64
Education                int64
EducationField           object
EmployeeCount            int64
EmployeeNumber           int64
EnvironmentSatisfaction  int64
Gender                   object
HourlyRate               int64
JobInvolvement           int64
JobLevel                 int64
JobRole                  object
JobSatisfaction          int64
MaritalStatus            object
MonthlyIncome            int64
MonthlyRate              int64
NumCompaniesWorked       int64
Over18                   object
OverTime                 object
PercentSalaryHike        int64
PerformanceRating        int64
RelationshipSatisfaction  int64
StandardHours            int64
StockOptionLevel         int64
TotalWorkingYears        int64
TrainingTimesLastYear    int64
WorkLifeBalance          int64
YearsAtCompany           int64
YearsInCurrentRole       int64
YearsSinceLastPromotion  int64
YearsWithCurrManager     int64
dtype: object
```

```
In [7]: categorical=data.select_dtypes('object')
print(categorical.columns)

Index(['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole',
      'MaritalStatus', 'Over18', 'OverTime'],
      dtype='object')
```

```
In [9]: numerical=data.select_dtypes('int64','float64')
print(numerical.columns)


Index(['Age', 'Attrition', 'DailyRate', 'DistanceFromHome', 'Education',
      'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfaction',
      'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobSatisfaction',
      'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
      'StandardHours', 'StockOptionLevel', 'TotalWorkingYears',
      'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
      'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager'],
      dtype='object')
```

**Distribution of the dataset :**

Generally one of the first few steps in exploring the data would be to have a rough idea of how the features are distributed with one another. To do so, I shall invoke the familiar kdeplot function from the Seaborn plotting library and this generates bivariate plots as follows:


```
In [10]: sns.kdeplot(data['Age'])
```

```
Out[10]: <Axes: xlabel='Age', ylabel='Density'>
```

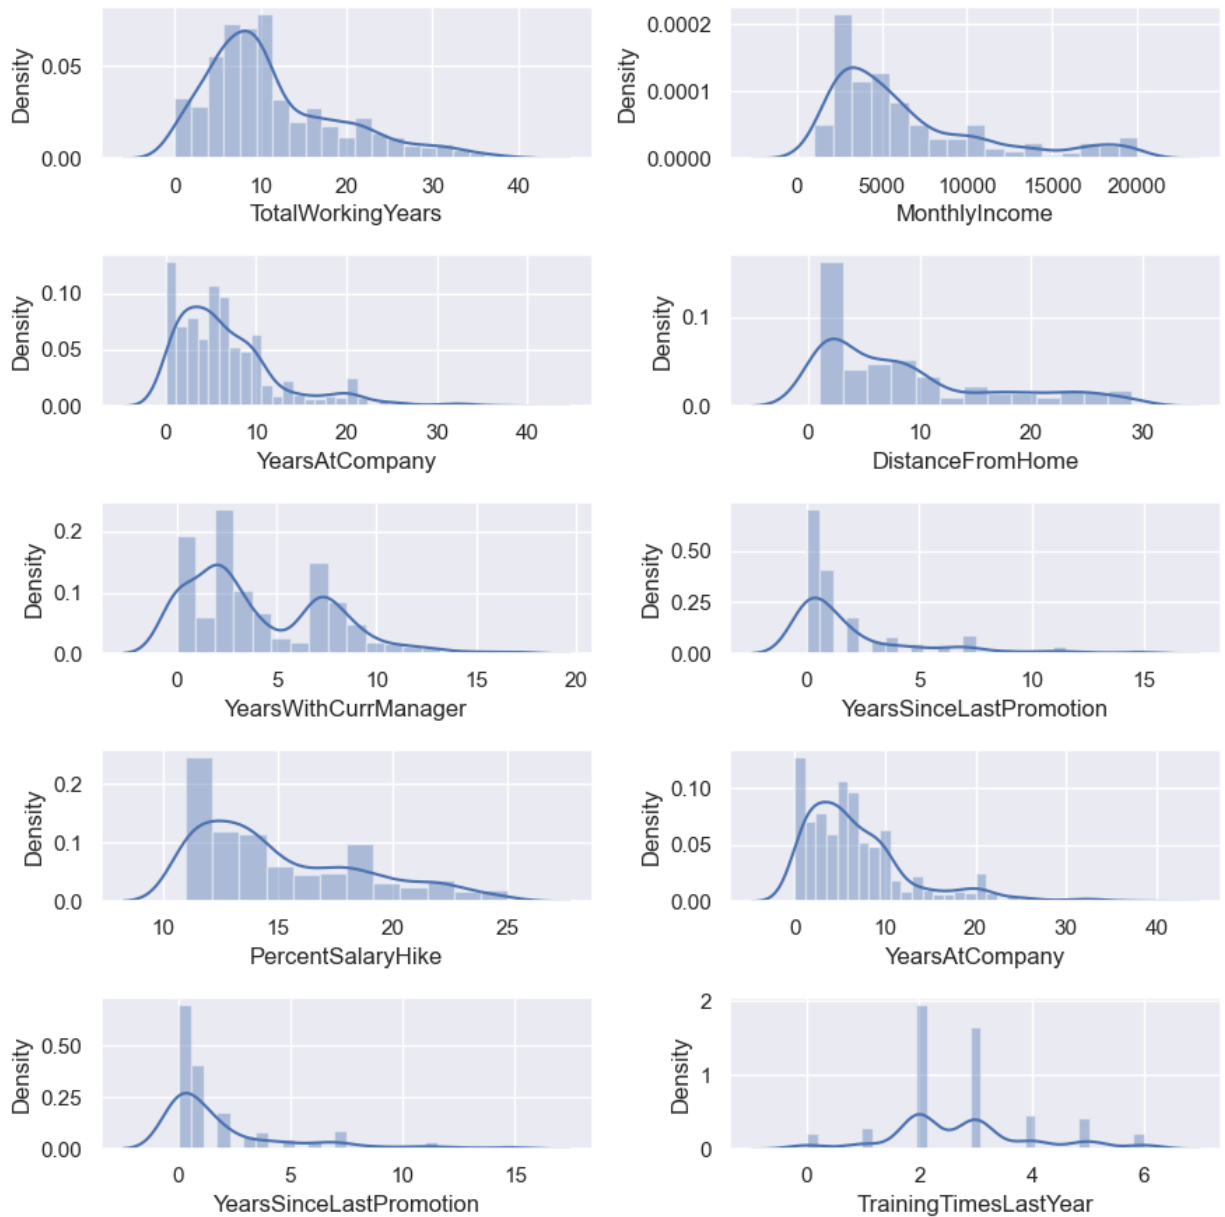


```
In [11]: sns.distplot(data['Age'])
```

```
Out[11]: <Axes: xlabel='Age', ylabel='Density'>
```



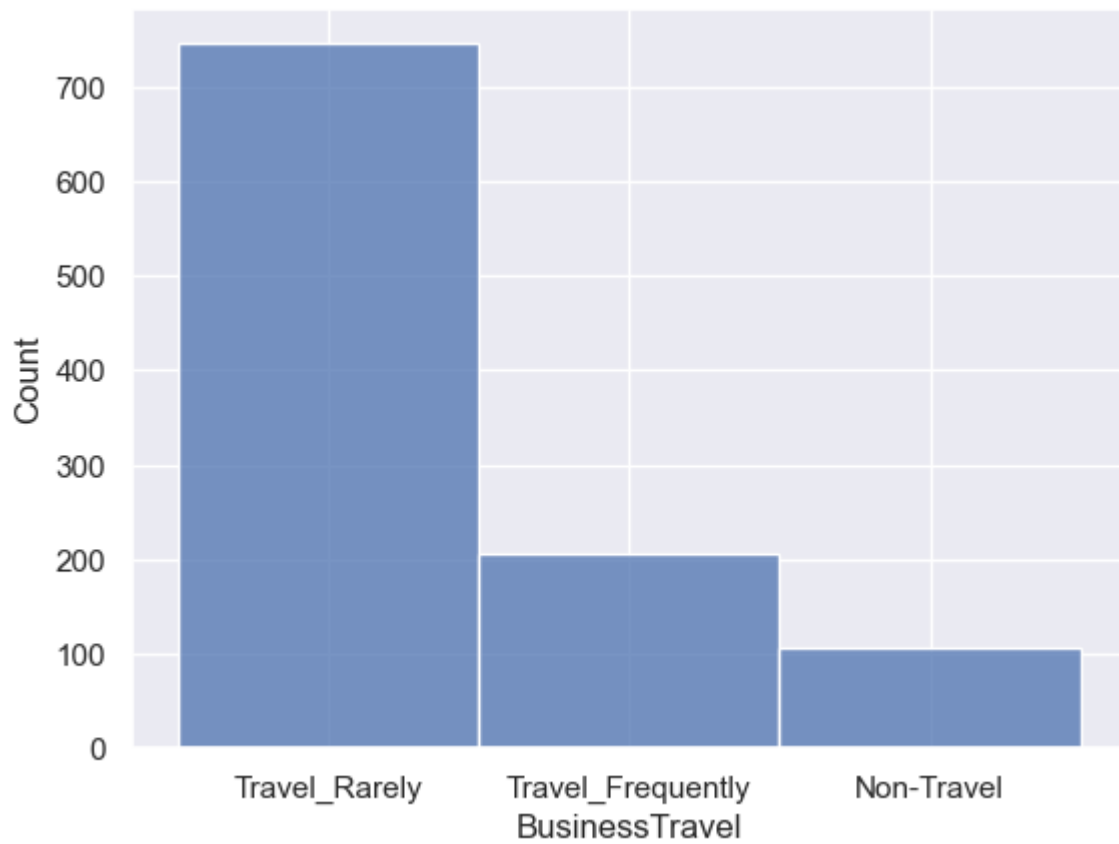
```
In [16]: fig, ax=plt.subplots(5,2,figsize=(9,9))
sns.distplot(data['TotalWorkingYears'],ax=ax[0,0])
sns.distplot(data['MonthlyIncome'],ax=ax[0,1])
sns.distplot(data['YearsAtCompany'],ax=ax[1,0])
sns.distplot(data['DistanceFromHome'],ax=ax[1,1])
sns.distplot(data['YearsWithCurrManager'],ax=ax[2,0])
sns.distplot(data['YearsSinceLastPromotion'],ax=ax[2,1])
sns.distplot(data['PercentSalaryHike'],ax=ax[3,0])
sns.distplot(data['YearsAtCompany'],ax=ax[3,1])
sns.distplot(data['YearsSinceLastPromotion'],ax=ax[4,0])
sns.distplot(data['TrainingTimesLastYear'],ax=ax[4,1])
plt.tight_layout()
plt.show()
```



```
In [17]: sns.histplot(data=data,x='BusinessTravel')
```

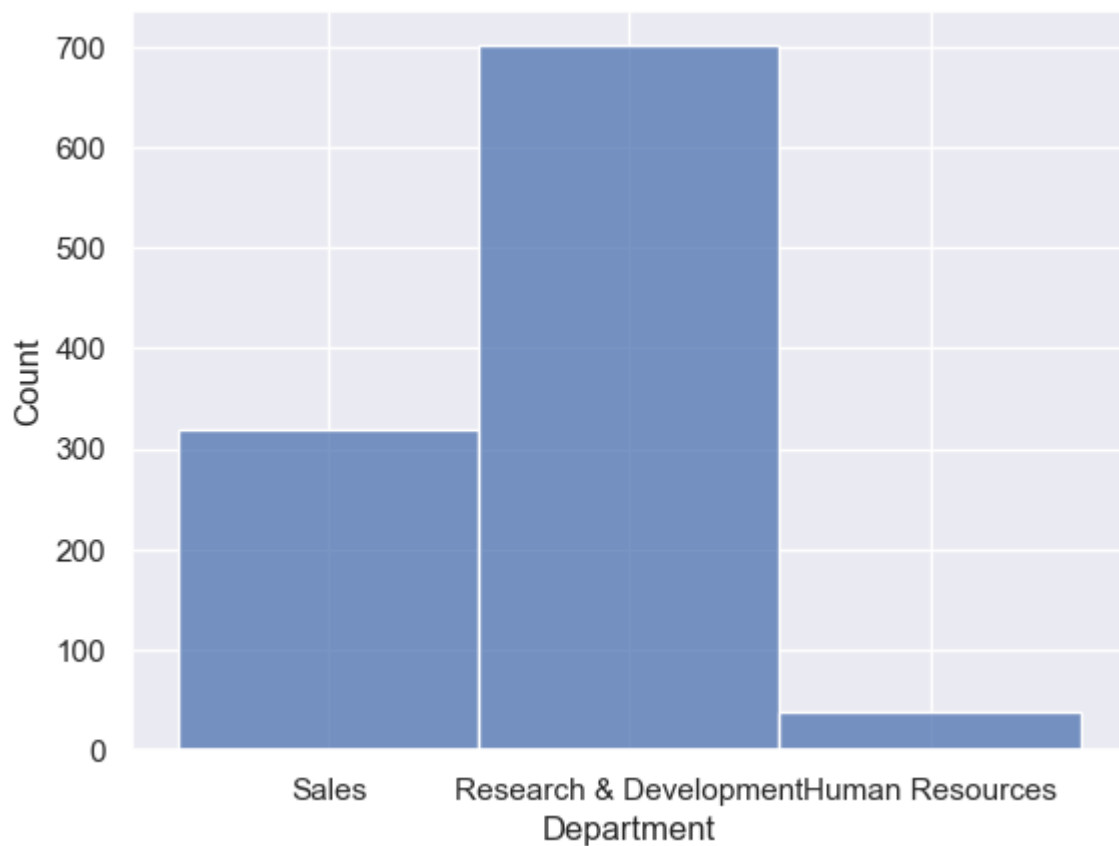
```
Out[17]: <Axes: xlabel='BusinessTravel', ylabel='Count'>
```






```
In [18]: sns.histplot(data=data, x='Department')
```

```
Out[18]: <Axes: xlabel='Department', ylabel='Count'>
```




```
In [19]: sns.histplot(data=data,y='EducationField')
```

```
Out[19]: <Axes: xlabel='Count', ylabel='EducationField'>
```




```
In [22]: bins=[0,18,35,np.inf]
labels=['Student','Freshers/Junior','Senior']
data['AgeGroup']=pd.cut(data['Age'],bins,labels=labels)
sns.histplot(data=data,y='AgeGroup')
```

```
Out[22]: <Axes: xlabel='Count', ylabel='AgeGroup'>
```



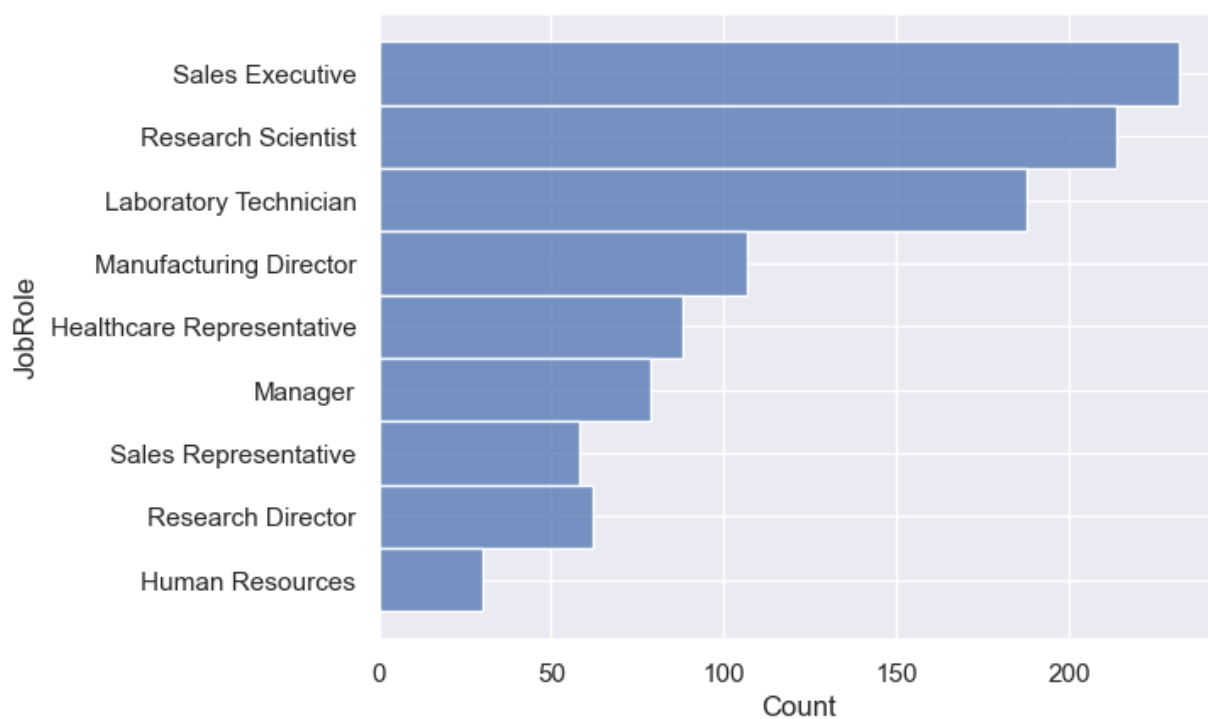
```
In [23]: sns.histplot(data=data,y='Gender')
```

```
Out[23]: <Axes: xlabel='Count', ylabel='Gender'>
```




```
In [24]: sns.histplot(data=data,y='JobRole')
```

Out[24]: <Axes: xlabel='Count', ylabel='JobRole'>




In [26]: `sns.histplot(data=data,x='Over18')`

Out[26]: <Axes: xlabel='Over18', ylabel='Count'>



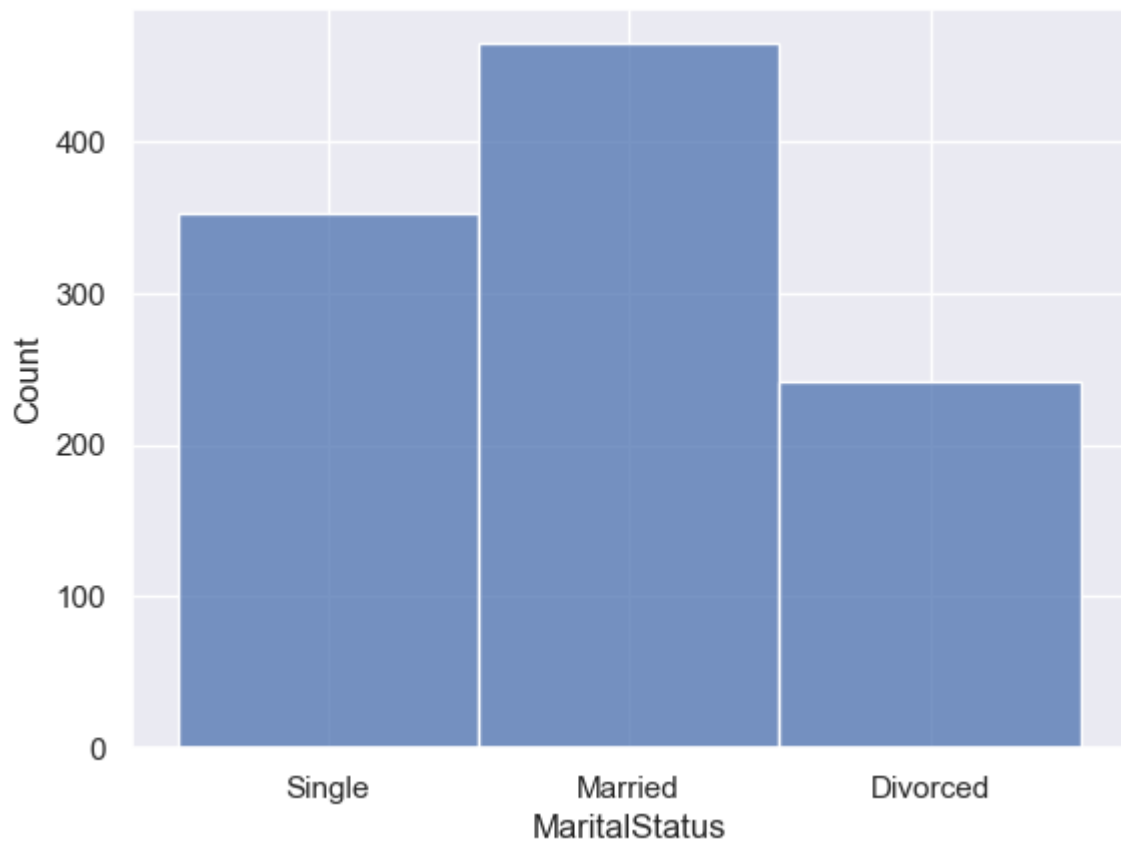
In [27]: `sns.histplot(data=data,x='OverTime')`

Out[27]: <Axes: xlabel='OverTime', ylabel='Count'>




In [29]: `sns.histplot(data=data,x='MaritalStatus')`

Out[29]: <Axes: xlabel='MaritalStatus', ylabel='Count'>



```
In [31]: sns.histplot(data=data,x='Attrition')
```

```
Out[31]: <Axes: xlabel='Attrition', ylabel='Count'>
```




## Correlation of Features

The next tool in a data explorer's arsenal is that of a correlation matrix. By plotting a correlation matrix, we have a very nice overview of how the features are related to one another. For a Pandas dataframe, we can conveniently use the call `.corr` which by default provides the Pearson Correlation values of the columns pairwise in that dataframe.

In this correlation plot, I will use the the Plotly library to produce a interactive Pearson correlation matrix via the Heatmap function as follows:

```
In [33]: cor_mat=data.corr()
mask=np.array(cor_mat)
mask[np.tril_indices_from(mask)]=False
fig=plt.gcf()
fig.set_size_inches(60,12)
sns.heatmap(data=cor_mat,mask=mask,square=True,annot=True,cbar=True,cmap="crest")
```

Out[33]: <Axes: >



```
In [34]: data.columns
```

```
Out[34]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
        'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
        'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
        'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
        'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
        'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
        'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
        'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
        'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
        'YearsWithCurrManager', 'AgeGroup'],
        dtype='object')
```


```
In [35]: continous=['Age','DailyRate','HourlyRate','MonthlyIncome','MonthlyRate','TotalWorking
```

```
In [42]: for var in continous:


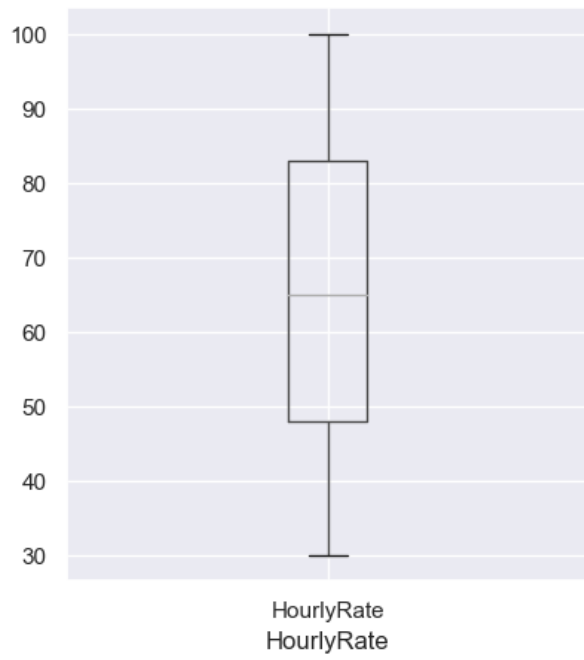


    # BoxPlot...
    plt.figure(figsize=(10,5))
    plt.subplot(1,2,1)
    fig=data.boxplot(column=var)
    fig.set_xlabel(var)



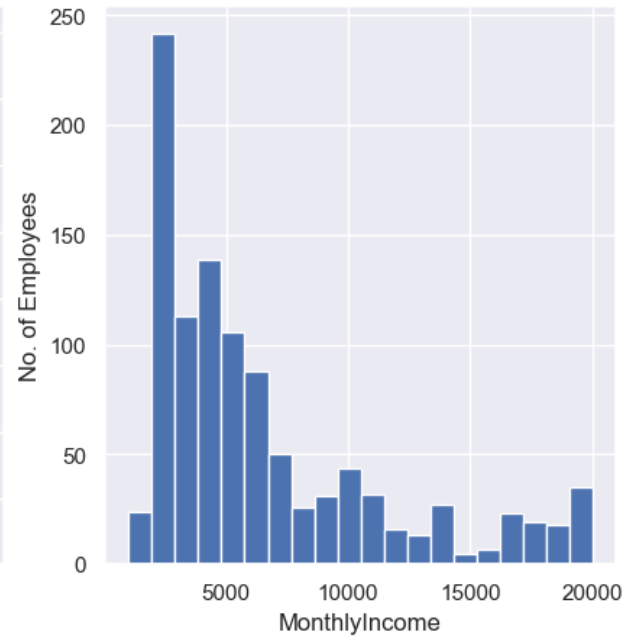
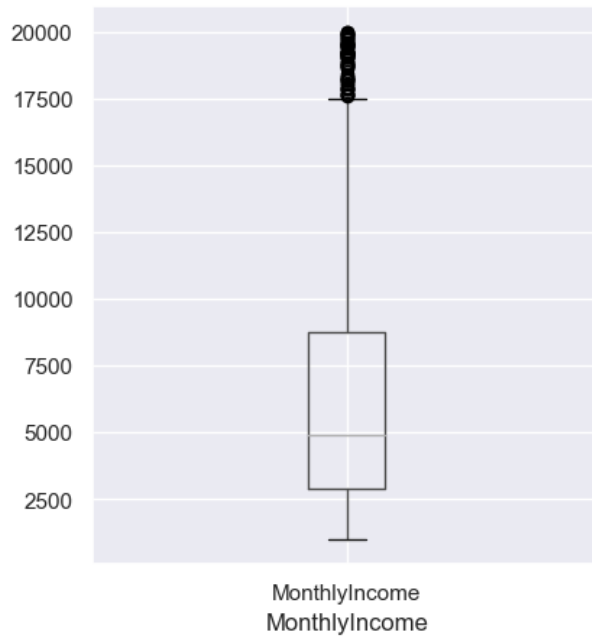
    # HistPlot...
    plt.subplot(1,2,2)
    fig=data[var].hist(bins=20)
    fig.set_ylabel('No. of Employees')
    fig.set_xlabel(var)





    plt.show()
```











```
In [43]: data['TotalWorkingYears'].describe()
```

```
Out[43]: count    1058.000000
mean       11.435728
std        8.016429
min         0.000000
25%        6.000000
50%       10.000000
75%       16.000000
max       40.000000
Name: TotalWorkingYears, dtype: float64
```

```
In [44]: categorical.head()
```

Out[44]:	BusinessTravel	Department	EducationField	Gender	JobRole	MaritalStatus	Over18	OverTime
0	Travel_Rarely	Sales	Life Sciences	Female	Sales Executive	Single	Y	Yes
1	Travel_Frequently	Research & Development	Life Sciences	Male	Research Scientist	Married	Y	No
2	Travel_Rarely	Research & Development	Other	Male	Laboratory Technician	Single	Y	Yes
3	Travel_Frequently	Research & Development	Life Sciences	Female	Research Scientist	Married	Y	Yes
4	Travel_Rarely	Research & Development	Medical	Male	Laboratory Technician	Married	Y	No

In [45]: `data_cat=pd.get_dummies(categorical)`

In [46]: `data_cat.head()`

Out[46]:	BusinessTravel_Non-Travel	BusinessTravel_Travel_Frequently	BusinessTravel_Travel_Rarely	Department_Human Resources
0	0	0	1	
1	0	1	0	
2	0	0	1	
3	0	1	0	
4	0	0	1	

5 rows × 29 columns

In [47]: `numerical.head()`

Out[47]:	Age	Attrition	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Environment
0	41	1	1102	1	2	1	1	
1	49	0	279	8	1	1	2	
2	37	1	1373	2	2	1	4	
3	33	0	1392	3	4	1	5	
4	27	0	591	2	1	1	7	

5 rows × 27 columns

In [49]: `data_final=pd.concat([numerical,data_cat], axis=1)`

In [50]: `data_final.head()`

Out[50]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Envir
0	41	1	1102	1	2	1	1	
1	49	0	279	8	1	1	2	
2	37	1	1373	2	2	1	4	
3	33	0	1392	3	4	1	5	
4	27	0	591	2	1	1	7	

5 rows × 56 columns

```
In [52]: data_final=data_final.drop('Attrition',axis=1)
data_final
```

Out[52]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Environmen
0	41	1102	1	2	1	1	
1	49	279	8	1	1	2	
2	37	1373	2	2	1	4	
3	33	1392	3	4	1	5	
4	27	591	2	1	1	7	
...	...	...	...	...	...	...	...
1053	57	405	1	2	1	1483	
1054	49	1490	7	4	1	1484	
1055	34	829	15	3	1	1485	
1056	28	1496	1	3	1	1486	
1057	29	115	13	3	1	1487	

1058 rows × 55 columns

```
In [56]: target=data['Attrition']
```

## Build Basline Models

```
In [58]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
```

```
In [60]: X_train,X_test,y_train,y_test = train_test_split(data_final , target ,test_size=0.2, r
```

```
In [61]: X_train.shape
```

```
Out[61]: (846, 55)
```

```
In [62]: X_test.shape
```

```
Out[62]: (212, 55)
```

```
In [65]: model=RandomForestClassifier()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
print('Accuracy : ', accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
```

```
Accuracy : 0.8915094339622641
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	182
1	0.89	0.27	0.41	30
accuracy			0.89	212
macro avg	0.89	0.63	0.68	212
weighted avg	0.89	0.89	0.87	212

```
In [67]: model=LogisticRegression()
model.fit(X_train,y_train)
y_predictions=model.predict(X_test)
print('Accuracy:', accuracy_score(y_test,y_predictions))
print(classification_report(y_test,y_predictions))
```

```
Accuracy: 0.8632075471698113
```

	precision	recall	f1-score	support
0	0.86	1.00	0.93	182
1	1.00	0.03	0.06	30
accuracy			0.86	212
macro avg	0.93	0.52	0.50	212
weighted avg	0.88	0.86	0.80	212

```
In [69]: model=DecisionTreeClassifier()
model.fit(X_train,y_train)
y_predictions = model.predict(X_test)
print('Accuracy : ', accuracy_score(y_test,y_predictions))
print(classification_report(y_test,y_predictions))
```

Accuracy :	0.7830188679245284				
	precision	recall	f1-score	support	
0	0.89	0.86	0.87	182	
1	0.28	0.33	0.30	30	
accuracy			0.78	212	
macro avg	0.58	0.60	0.59	212	
weighted avg	0.80	0.78	0.79	212	

```
In [71]: model = KNeighborsClassifier()
model.fit(X_train,y_train)
y_predictions = model.predict(X_test)
print('Accuracy', accuracy_score(y_test,y_predictions))
print(classification_report(y_test,y_predictions))
```

Accuracy	0.8254716981132075				
	precision	recall	f1-score	support	
0	0.86	0.95	0.90	182	
1	0.18	0.07	0.10	30	
accuracy			0.83	212	
macro avg	0.52	0.51	0.50	212	
weighted avg	0.76	0.83	0.79	212	

```
In [73]: model = SVC()
model.fit(X_train,y_train)
model_predictions = model.predict(X_test)
print('Accuracy : ', accuracy_score(y_test,y_predictions))
print(classification_report(y_test,y_predictions))
```

Accuracy :	0.8254716981132075				
	precision	recall	f1-score	support	
0	0.86	0.95	0.90	182	
1	0.18	0.07	0.10	30	
accuracy			0.83	212	
macro avg	0.52	0.51	0.50	212	
weighted avg	0.76	0.83	0.79	212	

## Insights

**In above I use some Machine Learning Algorithm to find the Accuracy Score of Test data and Predict Data:**

1. First I use a Train-Test Split to evaluate a machine learning model's ability to predict a certain outcome accurately when exposed to real-world data it's never seen before. In a train-test split, I split an original dataset into two subsets—a training dataset and a testing dataset. Depending on the nature and complexity of the data. Then find a Predict data and move to next step.

1. Next I use a Five Machine Learning Algorithm i.e. :

- "Random Forest Classifier" where Accuracy is 0.8915094339622641
- "Logistic Regression" where Accuracy is 0.8632075471698113
- "Decision Tree Classifier" where Accuracy is 0.7830188679245284
- "K Neighbours Classifier" where Accuracy is 0.8254716981132075
- "SVC" where Accuracy is 0.8254716981132075

In [ ]: